



relu: class structure by depth, PCA of hidden activations


gradient descent | layer 1 (23% var)




gradient descent | layer 4 (24% var)




gradient descent | layer 8 (81% var)




Adam | layer 1 (20% var)




Adam | layer 4 (65% var)




Adam | layer 8 (90% var)



reference (operator) | layer 1 (23% var)



reference (operator) | layer 4 (34% var)



reference (operator) | layer 8 (97% var)

