



relu: class structure by depth, PCA of hidden activations


gradient descent | layer 1 (23% var)




gradient descent | layer 4 (24% var)




gradient descent | layer 8 (81% var)




Adam | layer 1 (21% var)




Adam | layer 4 (44% var)




Adam | layer 8 (77% var)



reference (operator) | layer 1 (28% var)



reference (operator) | layer 4 (39% var)



reference (operator) | layer 8 (100% var)

